

Deep Learning Bootcamp with PyTorch

Noel Shin

Contents

- **Introduction**
- **Basics**
- **Appendix**

Introduction

Goal of lecture

Getting used to


- read a PyTorch code of interest,
- make a deep learning model using PyTorch.

What is PyTorch?

PyTorch is an **open-source machine learning library** for Python, based on Torch.

- **Tensor computation with strong GPU acceleration***
- Deep neural networks built on a tape-based **autograd system**

Why PyTorch?




Keras

PYTORCH

- A lot easier to learn than to learn TensorFlow.
- Available to customize a model. This is usually infeasible with Keras.
- Public codes for deep learning research papers are usually written with PyTorch.



Basics

What is deep learning?



Classification Model


How does it work?



Penguin: 0.01

Iteration

Difference
measurement
(Loss)

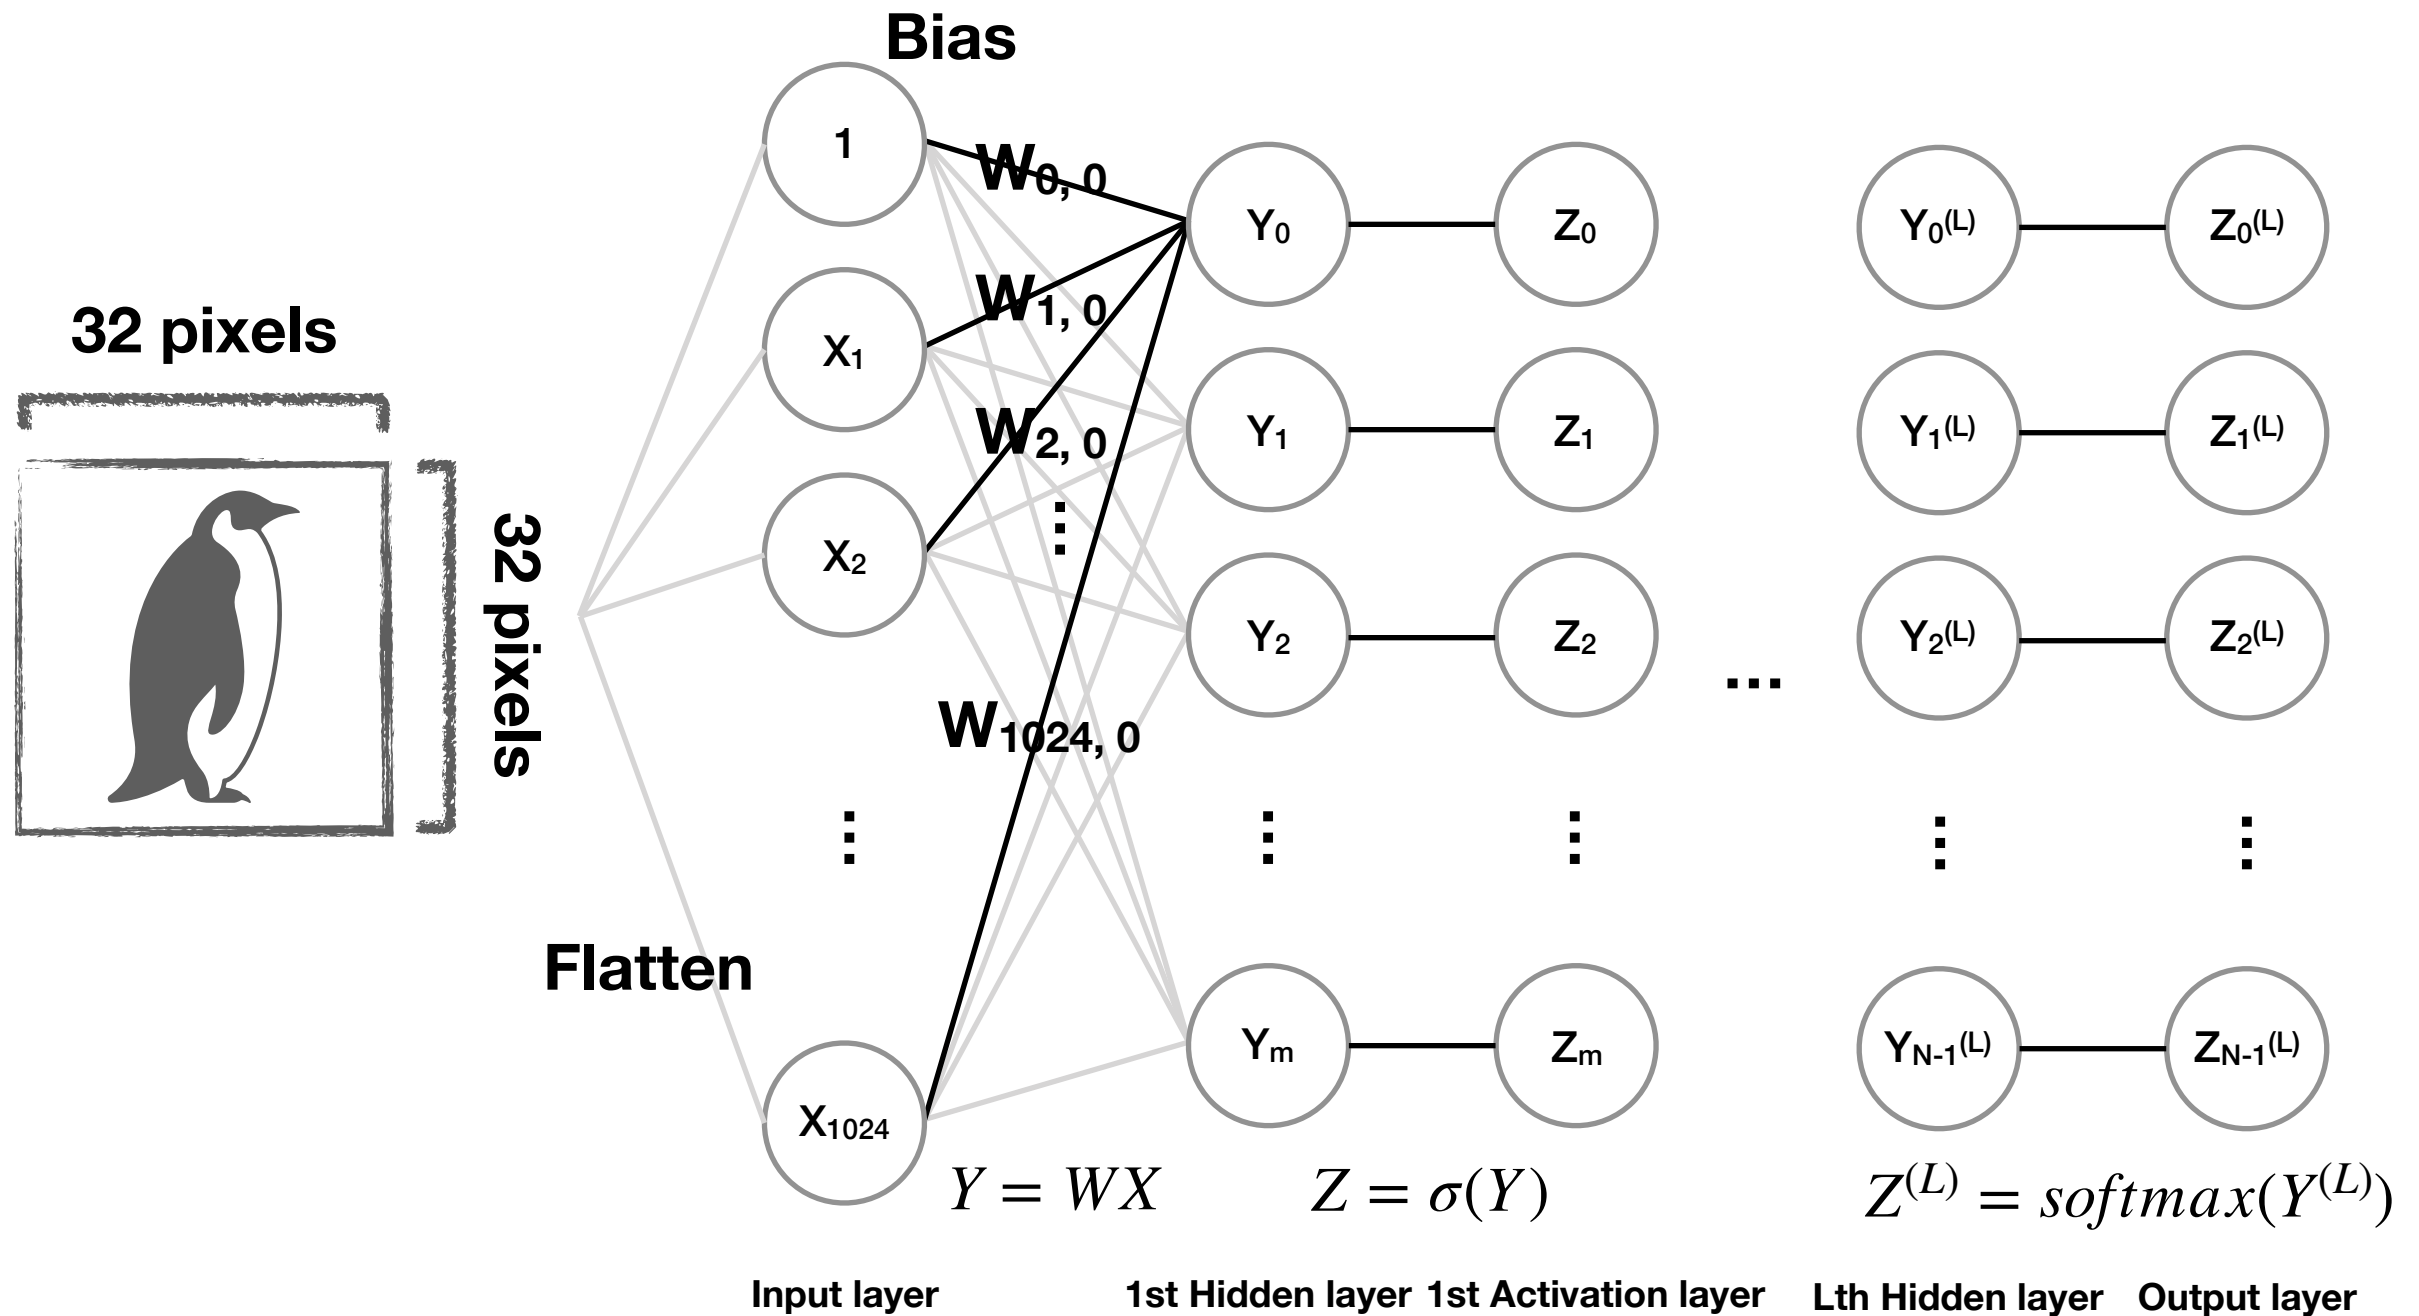


Backward path
(Gradient Back-propagation)

$\nabla_{\theta} L$

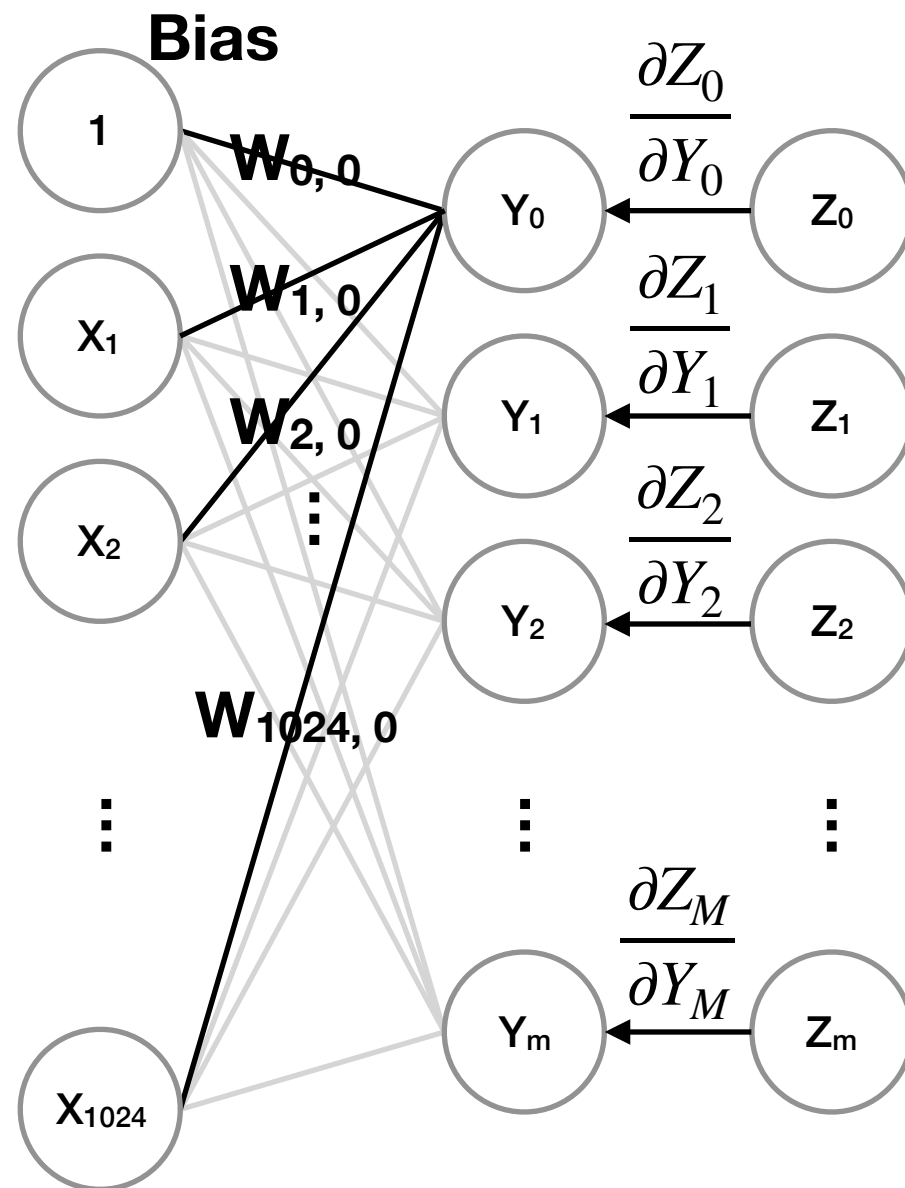
Penguin: 1.0

Forward path



**Densely connected
(fully connected)**


Backward path



$$\frac{\partial L}{\partial W_{0,0}} = \frac{\partial L}{\partial Y_0} \frac{\partial Y_0}{\partial W_{0,0}}$$

$$W_{0,0} = W_{0,0} - \rho \frac{\partial L}{\partial W_{0,0}}$$


...



$$L = - \sum_{k=0}^{N-1} t_k \log Z_k = - \log Z_i$$


Cross-entropy loss

Convolution*



H: Height
W: Width
OC: Output Channel
IC: Input Channel
KH: Kernel Height
KW: Kernel Width

Convolutional neural network



CONV: Convolution
BN: Batch Normalization
RELU: REctified Linear Unit
FC: Fully Connected layer
SM: SoftMax activation layer

Practice

MNIST database



Training set: 60,000 images and labels
Test set: 10,000 images and labels

The MNIST database (Modified National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.


Generative Model

Various generative models

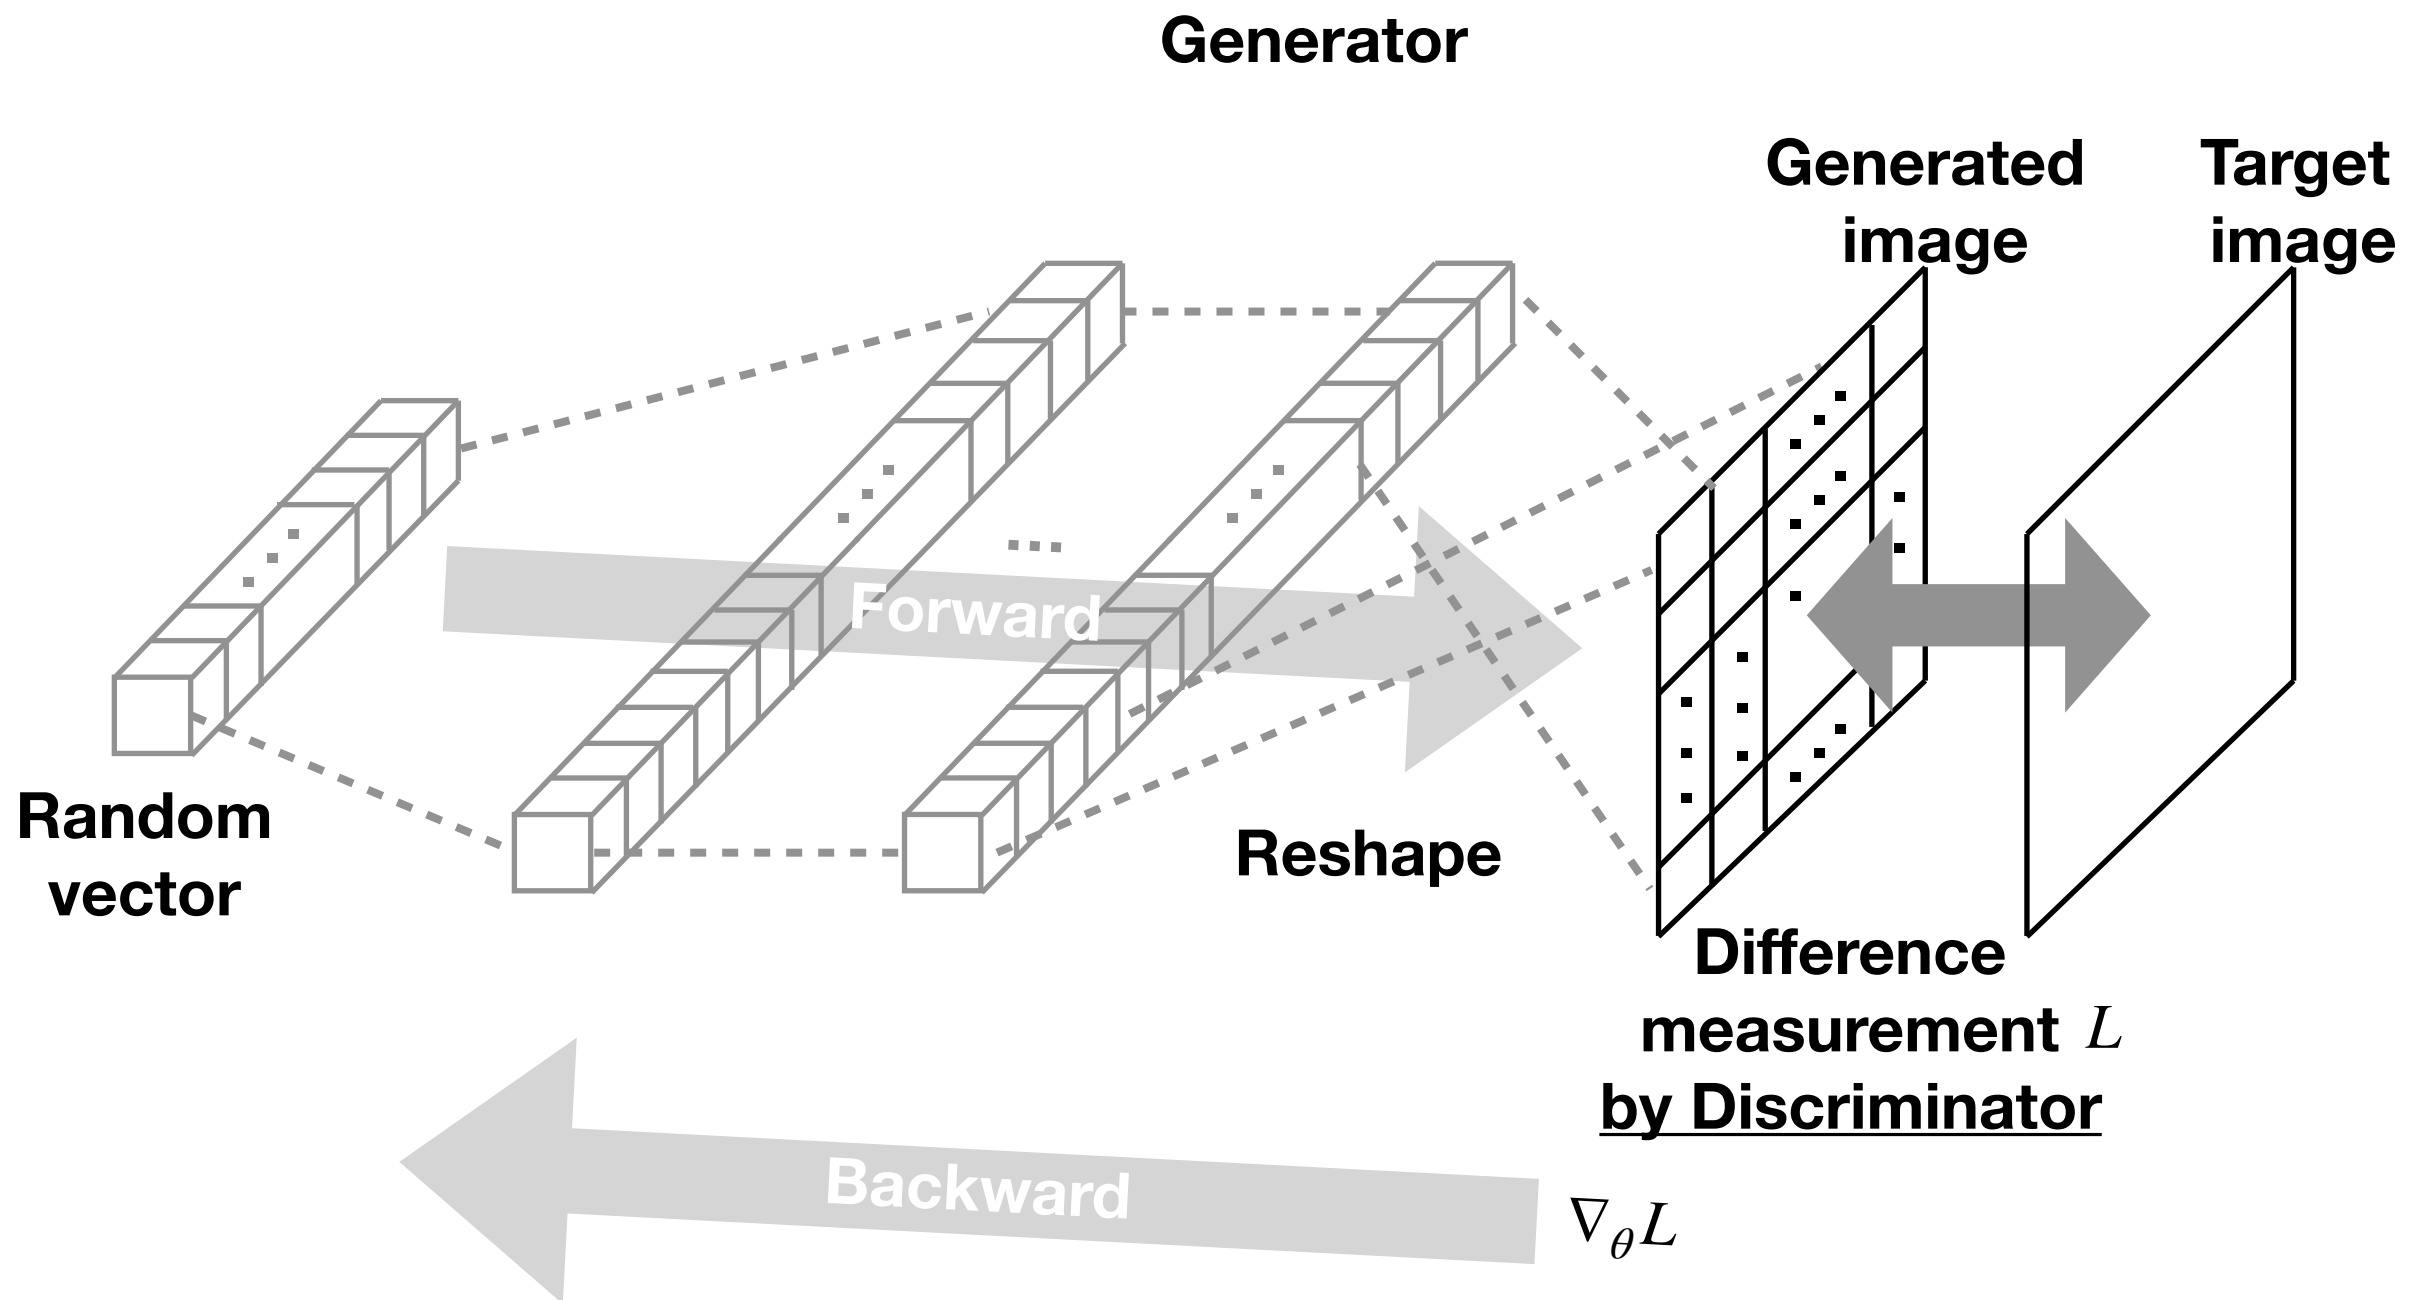
- Hidden Markov Model (HMM)
- Restricted Boltzmann Machine (RBM)
- Variational Auto-Encoder (VAE)
- Recurrent Neural Network (RNN)
- **Generative Adversarial Network (GAN)**

GAN

What is GAN?

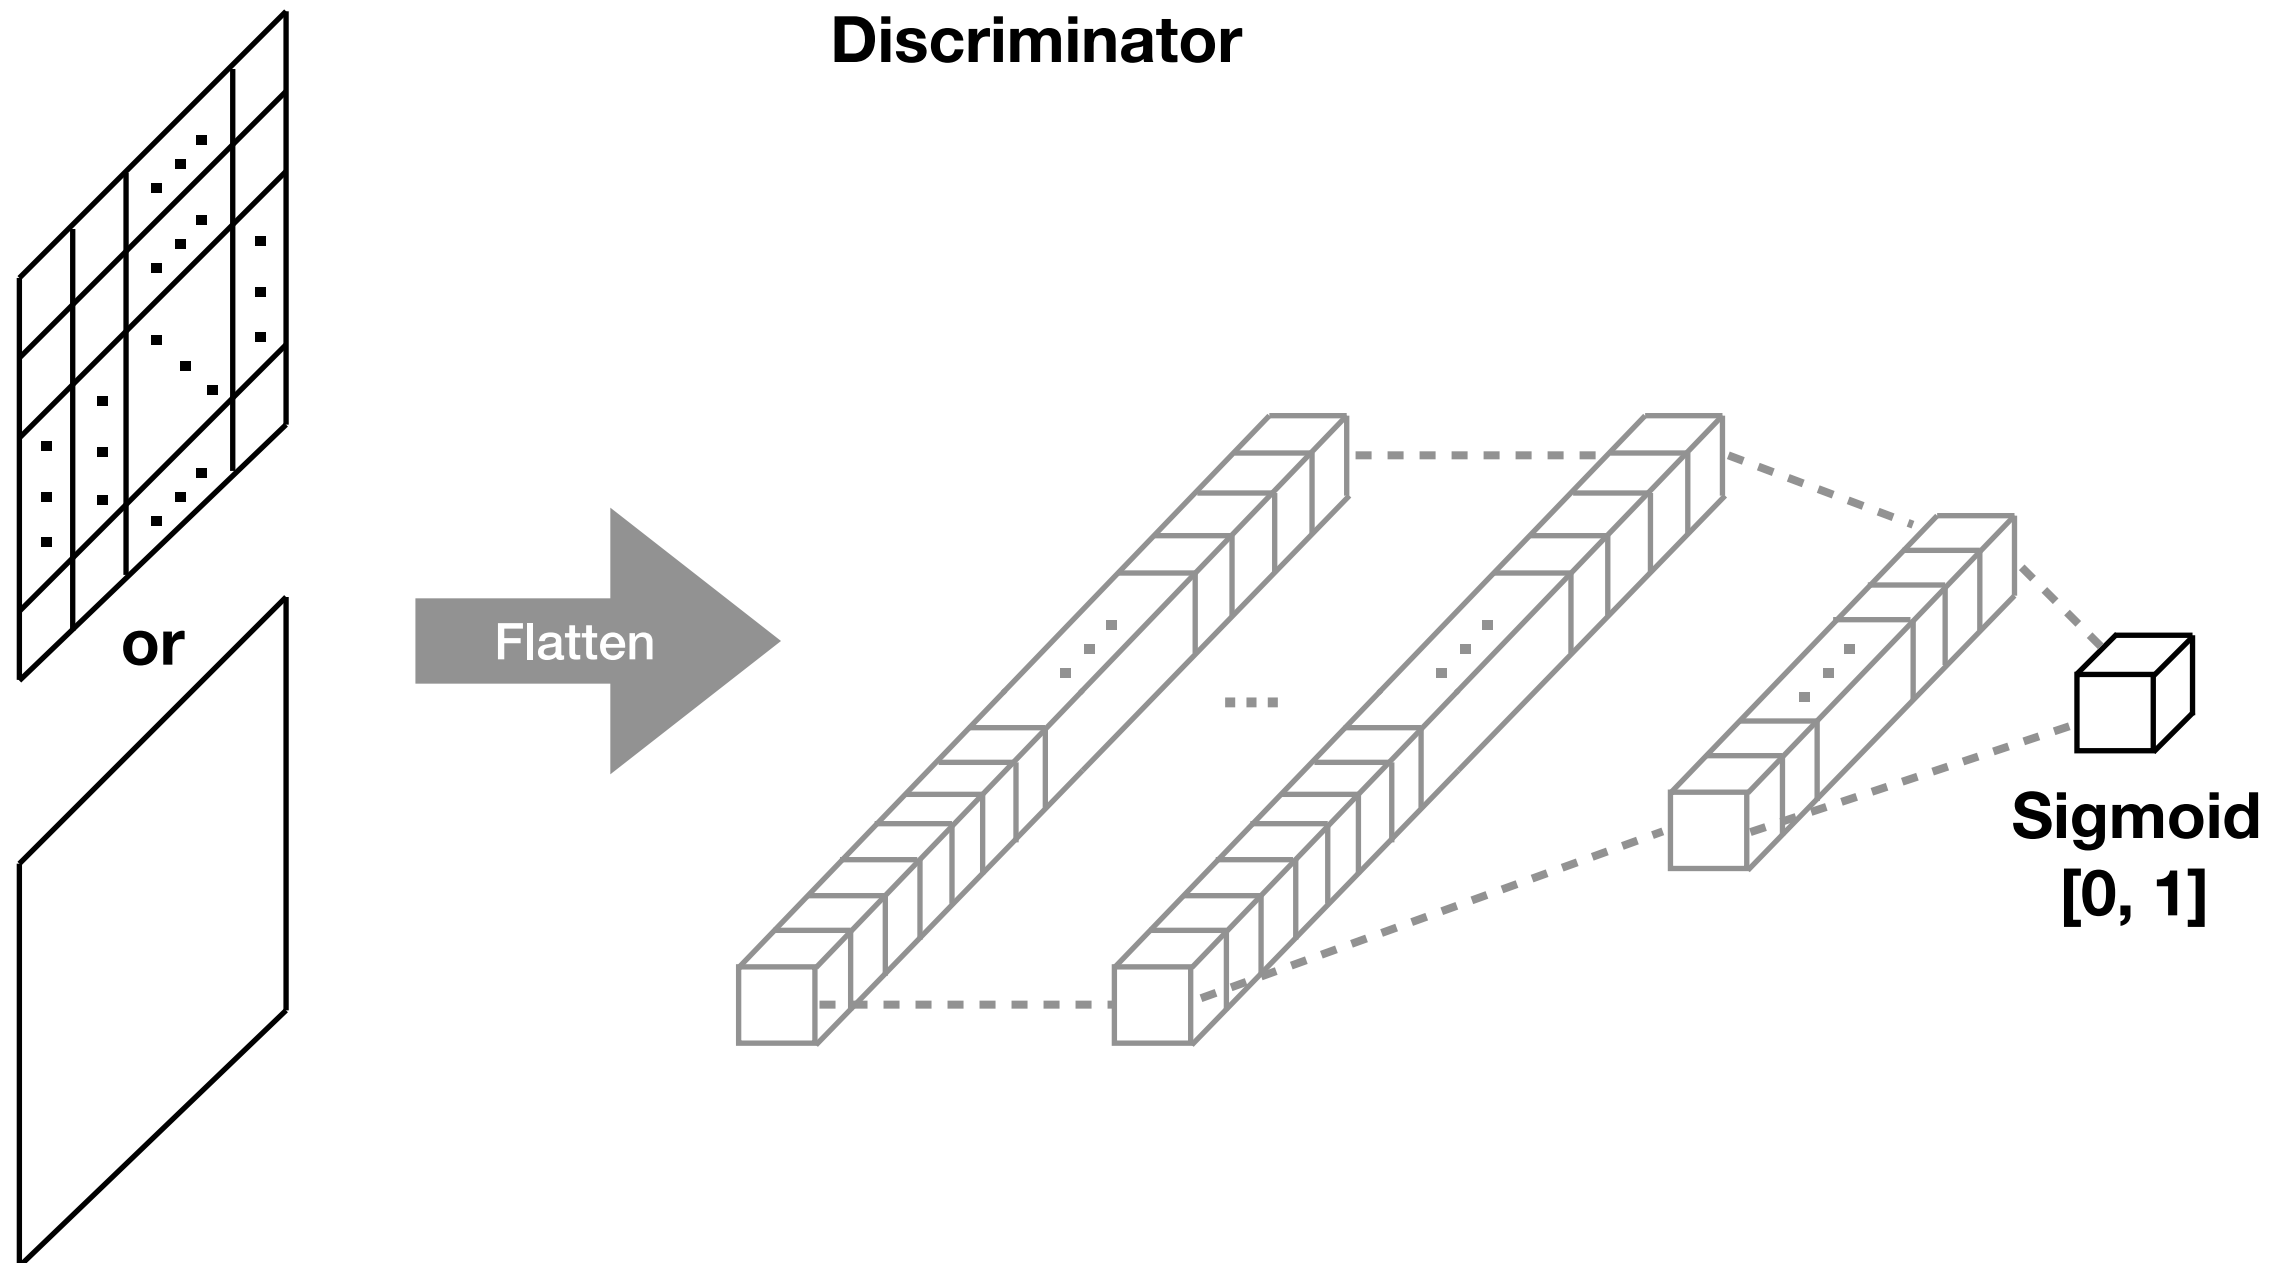


How does GAN work?



E.g. image generation model

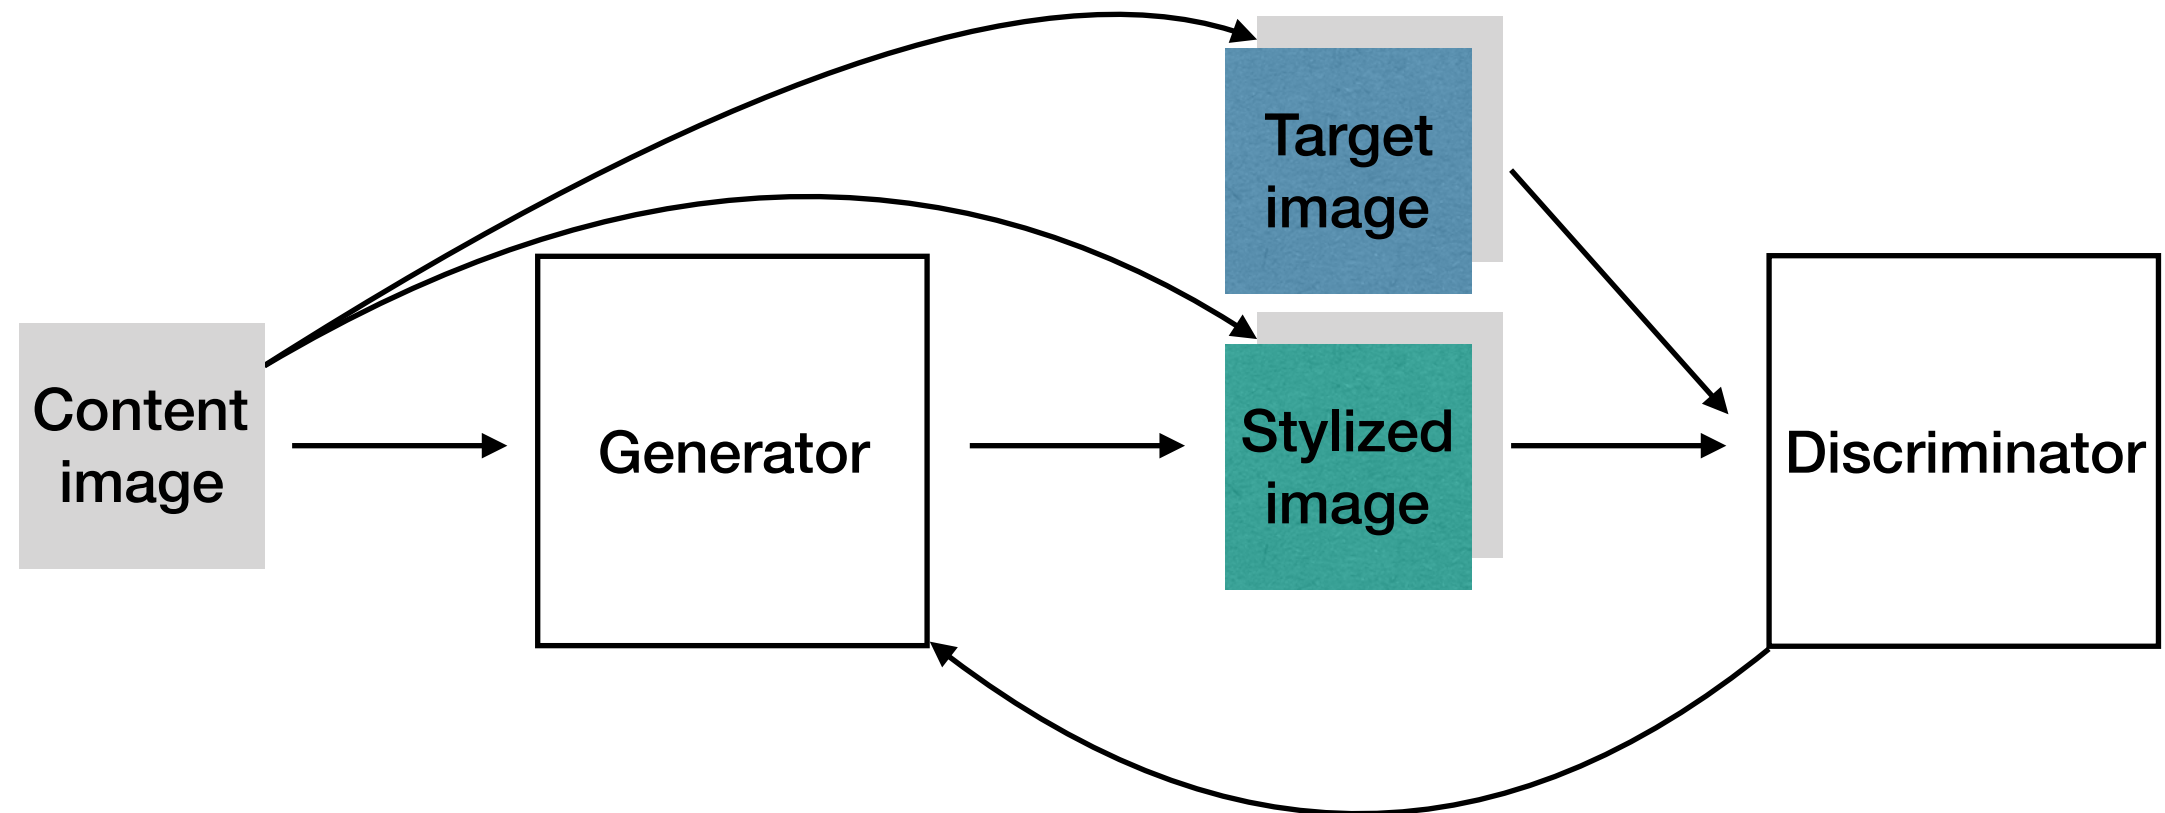
How does GAN work?



Practice

Conditional GAN

What is cGAN?



Practice


Appendix

CPU vs. GPU



credit. [hothardware.com](https://www.hothardware.com)





A central processing unit (CPU), also called a central processor or main processor, is the electronic circuitry within a computer that carries out the instructions of a computer program by performing the basic arithmetic, logic, controlling, and input/output operations specified by the instructions.



credit. NVIDIA

A graphics processing unit (GPU) is a specialized electronic circuit designed to rapidly manipulate and alter memory to accelerate the creation of images in a frame buffer intended for output to a display device.

Activation functions




- ReLU (Rectified linear unit) : $\max(0, x)$
- Leaky ReLU : $\max(0, x) + \text{negative slope} \times \min(0, x)$
- Hyperbolic tangent (tanh): $\frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Logistic sigmoid : $\frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1} = \frac{1}{2} + \frac{1}{2}\tanh\left(\frac{x}{2}\right)$

Loss functions


- Binary cross entropy loss (BCE) : $-\log \hat{p}(y_c | x) - \sum_{k=1, k \neq c}^C \log[1 - \hat{p}(y_k | x)]$
- Cross entropy loss (CE) : $\mathbb{E}_{p(y|x)}[-\log \hat{p}(y | x)] = - \sum_{c=1}^C p(y_c | x) \log \hat{p}(y_c | x) = - \log \hat{p}(y_c | x)$
- Mean squared error loss (MSE) : $\frac{1}{MN} \sum_{i=0}^{M-1} \sum_{j=0}^{N-1} (x_{i,j} - \hat{x}_{i,j})^2$
- Mean absolute error loss (MAE) : $\frac{1}{MN} \sum_{i=0}^{M-1} \sum_{j=0}^{N-1} |x_{i,j} - \hat{x}_{i,j}|$

Convolution




$$o = i - k + 1$$

Convolution




$(7 + 2 \times p) \times (7 + 2 \times p)$
Input

\otimes
Padding=1
Stride=1



3x3
Kernel


=



7x7
Output


$$o = i - k + 2p + 1$$

Convolution




$(7 + 2 \times p) \times (7 + 2 \times p)$
Input

\otimes
Padding=1
Stride=2



3x3
Kernel

=



4x4
Output

$$o = \left\lfloor \frac{i - k + 2p}{s} \right\rfloor + 1$$